
Course Code: CPE 019Code Title: Emerging Technologies in CpE 2 - 2nd Semester

ACTIVITY NO.**Assignment 8.1 : Saving Models****Name**

Dela Cruz, Irish

Section

CPE32S3

Date Performed:

04/10/2024

Date Submitted:

04/19/2024

Instructor:Engr. Roman M. Richard

Introduction of Data

The "Census Income" dataset, a.k.a the "Adult" dataset, is a widely used dataset in the machine learning community, particularly for classification tasks. It contains demographic information from the US Census Bureau and is commonly used to predict whether an individual's income exceeds \$50K per year based on various attributes

Link: <https://archive.ics.uci.edu/dataset/2/adult>

Dataset:

- Multivariate Subject Area:
- Social Science

Features: (categorical and integer type)

- age
- education level
- marital status
- occupation
- race
- gender
- capital gain
- capital loss
- hours worked per week
- native country

Instances:

- 48,842

This dataset is mostly used for task such as income prediction, demographic analysis, understanding the factors that influence income levels, and eploring data pre-processing techniques due to its real-world relevance and moderate size.

➤ Installing the requirement packages

[] ↳ 2 cells hidden

✓ Explonatory Data Analysis (EDA)

The following task were performed in this section:

- 1. Checking the data is properly loaded using head and tail
- 2. Checking the information of the dataframes
- 3. Checking the features by graphing all the values
- 4. Checking the values count of each variables

X.shape, y.shape


((48842, 14), (48842, 1))

X.head(10)

	age	workclass	fnlwgt	education	educati
0	39	State-gov	77516	Bachelors	
1	50	Self-emp-not-inc	83311	Bachelors	
2	38	Private	215646	HS-grad	
3	53	Private	234721	11th	
4	28	Private	338409	Bachelors	

Next steps: [View recommended plots](#)

```
y.head(10)
```

	income	
0	<=50K	
1	<=50K	
2	<=50K	
3	<=50K	
4	<=50K	
5	<=50K	
6	<=50K	
7	>50K	
8	>50K	
9	>50K	

Next steps:

 [View recommended plots](#)

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   48842 non-null  int64
1   workclass             47879 non-null  object
2   fnlwgt                48842 non-null  int64
3   education             48842 non-null  object
4   education-num         48842 non-null  int64
5   marital-status        48842 non-null  object
6   occupation            47876 non-null  object
7   relationship          48842 non-null  object
8   race                  48842 non-null  object
9   sex                   48842 non-null  object
10  capital-gain           48842 non-null  int64
11  capital-loss           48842 non-null  int64
12  hours-per-week         48842 non-null  int64
13  native-country        48568 non-null  object
dtypes: int64(6), object(8)
memory usage: 5.2+ MB
```

```
y.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype

```

```
--- -----
0    income  48842 non-null  object
dtypes: object(1)
memory usage: 381.7+ KB
```

✓ Separating Categorical and Numerical Variables

```
numerical_features = X.select_dtypes(include=['float64', 'int64'])
categorical_features = X.select_dtypes(include=['object'])

print("Numerical Features:")
print(numerical_features.columns)

print("\nCategorical Features:")
print(categorical_features.columns)

Numerical Features:
Index(['age', 'fnlwgt', 'education-num', 'capital-gain', 'capital-loss',
       'hours-per-week'],
      dtype='object')

Categorical Features:
Index(['workclass', 'education', 'marital-status', 'occupation',
       'relationship', 'race', 'sex', 'native-country'],
      dtype='object')
```

✓ Handling Categorical Variable

```
from sklearn import preprocessing
label_encoder = preprocessing.LabelEncoder()
X['workclass'] = label_encoder.fit_transform(X['workclass'])
X['education'] = label_encoder.fit_transform(X['education'])
X['marital-status'] = label_encoder.fit_transform(X['marital-status'])
X['occupation'] = label_encoder.fit_transform(X['occupation'])
X['relationship'] = label_encoder.fit_transform(X['relationship'])
X['race'] = label_encoder.fit_transform(X['race'])
X['sex'] = label_encoder.fit_transform(X['sex'])
X['native-country'] = label_encoder.fit_transform(X['native-country'])
X.head(20)
```

```
<ipython-input-11-9f5d11bed876>:3: SettingWith  
A value is trying to be set on a copy of a slice  
Try using .loc[row_indexer,col_indexer] = value
```

See the caveats in the documentation: <https://>

```
X['workclass'] = label_encoder.fit_transform  
<ipython-input-11-9f5d11bed876>:4: SettingWith  
A value is trying to be set on a copy of a slice  
Try using .loc[row_indexer,col_indexer] = value
```

See the caveats in the documentation: <https://>

```
X['education'] = label_encoder.fit_transform  
<ipython-input-11-9f5d11bed876>:5: SettingWith  
A value is trying to be set on a copy of a slice  
Try using .loc[row_indexer,col_indexer] = value
```

See the caveats in the documentation: <https://>

```
X['marital-status'] = label_encoder.fit_trans  
<ipython-input-11-9f5d11bed876>:6: SettingWith  
A value is trying to be set on a copy of a slice  
Try using .loc[row_indexer,col_indexer] = value
```

See the caveats in the documentation: <https://>

```
X['occupation'] = label_encoder.fit_transfo  
<ipython-input-11-9f5d11bed876>:7: SettingWith  
A value is trying to be set on a copy of a slice  
Try using .loc[row_indexer,col_indexer] = value
```

See the caveats in the documentation: <https://>

```
X['relationship'] = label_encoder.fit_trans  
<ipython-input-11-9f5d11bed876>:8: SettingWith  
A value is trying to be set on a copy of a slice  
Try using .loc[row_indexer,col_indexer] = value
```

See the caveats in the documentation: <https://>

```
X['race'] = label_encoder.fit_transform(X['  
<ipython-input-11-9f5d11bed876>:9: SettingWith  
A value is trying to be set on a copy of a slice  
Try using .loc[row_indexer,col_indexer] = value
```

See the caveats in the documentation: <https://>

```
X['sex'] = label_encoder.fit_transform(X['s  
<ipython-input-11-9f5d11bed876>:10: SettingWi  
A value is trying to be set on a copy of a slice  
Try using .loc[row_indexer,col_indexer] = value
```

See the caveats in the documentation: <https://>

```
X['native-country'] = label_encoder.fit_tra
```

	age	workclass	fnlwgt	education	educat
0	39	7	77516	9	
1	50	6	83311	9	
2	38	4	215646	11	

3	53	4	234721	1
4	28	4	338409	9
5	37	4	284582	12
6	49	4	160187	6
7	52	6	209642	11
8	31	4	45781	12
9	42	4	159449	9
10	37	4	280464	15
11	30	7	141297	9
12	23	4	122272	9
13	32	4	205019	7
14	40	4	121772	8
15	34	4	245487	5
16	25	6	176756	11
17	32	4	186824	11

18 28 4 28887 1

Next steps:  [View recommended plots](#)

```
from sklearn import preprocessing
label_encoder = preprocessing.LabelEncoder()
y['income'] = label_encoder.fit_transform(y['income'])
```

```
<ipython-input-12-a399ca7721c1>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-y>
y['income'] = label_encoder.fit_transform(y['income'])

Remarks:

There's 9 categorical values which are 'workclass', 'education', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'native-country' and 'income'. I used label encoder to convert it into string since neural networks is crucial for enabling models to effectively learn from non-numerical data, improve performance, and avoid bias. So proper encoding and representation of categorical variables allow

neural networks to leverage the full potential of the available data and make more accurate predictions

✓ Verifying the conversion of Categorical into Numerical Features

X.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                   48842 non-null  int64
1   workclass             48842 non-null  int64
2   fnlwgt               48842 non-null  int64
3   education             48842 non-null  int64
4   education-num        48842 non-null  int64
5   marital-status       48842 non-null  int64
6   occupation           48842 non-null  int64
7   relationship         48842 non-null  int64
8   race                 48842 non-null  int64
9   sex                  48842 non-null  int64
10  capital-gain         48842 non-null  int64
11  capital-loss         48842 non-null  int64
12  hours-per-week       48842 non-null  int64
13  native-country       48842 non-null  int64
dtypes: int64(14)
memory usage: 5.2 MB
```

y.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  ---
0   income  48842 non-null  int64
dtypes: int64(1)
memory usage: 381.7 KB
```

✓ Handling Missing Values

X.isnull().sum()



```
age           0
workclass     0
fnlwgt       0
education     0
```

```

education-num    0
marital-status   0
occupation       0
relationship     0
race             0
sex             0
capital-gain     0
capital-loss     0
hours-per-week   0
native-country   0
dtype: int64

```

```
y.isnull()
```

	income	
0	False	
1	False	
2	False	
3	False	
4	False	
...	...	
48837	False	
48838	False	
48839	False	
48840	False	
48841	False	
48842 rows × 1 columns		

```

X["workclass"].fillna(X["workclass"].mean(), inplace=True)
X["occupation"].fillna(X["occupation"].mean(), inplace=True)
X["native-country"].fillna(X["native-country"].mean(), inplace=True)

```

```

<ipython-input-17-bcb232619632>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-c>

```

X["workclass"].fillna(X["workclass"].mean(), inplace=True)

```

```

<ipython-input-17-bcb232619632>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-c>

```

X["occupation"].fillna(X["occupation"].mean(), inplace=True)

```

```

<ipython-input-17-bcb232619632>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

```


See the caveats in the documentation: <https://pandas.pydata.org/pandas-c>
`X["native-country"].fillna(X["native-country"].mean(), inplace=True)`

✓ Verifying the missing values

```
X.isnull().sum()
```


```
age                0
workclass          0
fnlwgt            0
education         0
education-num     0
marital-status    0
occupation        0
relationship      0
race              0
sex              0
capital-gain      0
capital-loss      0
hours-per-week    0
native-country    0
dtype: int64
```

✓ Remarks for missing values

After the label being done converted the categorical transform into numerical datatype. All of the variables there are now ready to handle any pattern or model since there all in numerical form

```
X.describe()
```

	age	workclass	fnlwgt
count	48842.000000	48842.000000	4.884200e+
mean	38.643585	4.047889	1.896641e+
std	13.710510	1.528374	1.056040e+
min	17.000000	0.000000	1.228500e+
25%	28.000000	4.000000	1.175505e+
50%	37.000000	4.000000	1.781445e+
75%	48.000000	4.000000	2.376420e+



Remarks for describe function

The result above, regarding with the 25% and 75% is a big gap compared to each of them respectively. This proves that there's an outliers goes within the variables, to visualized and to verify those things a box plot will be performed

✓ Handling Outliers

```
X.columns, y.columns
```

```
(Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',  
        'marital-status', 'occupation', 'relationship', 'race', 'sex',  
        'capital-gain', 'capital-loss', 'hours-per-week', 'native-  
country'],  
      dtype='object'),  
 Index(['income'], dtype='object'))
```


```
import seaborn as sns  
import matplotlib.pyplot as plt  
import numpy as np  
import pandas as pd
```

```
columns = ['age', 'workclass', 'fnlwgt', 'education', 'education-num',  
           'marital-status', 'occupation', 'relationship', 'race', 'sex',  
           'capital-gain', 'capital-loss', 'hours-per-week', 'native-country']
```

```
fig, axs = plt.subplots(1, len(columns), figsize=(30, 8))
```

```
for i, col in enumerate(columns):  
    value_counts = X[col].value_counts()  
    axs[i].bar(value_counts.index, value_counts.values, color='lightcoral')  
    axs[i].set_xlabel(col)
```

```
plt.tight_layout()  
plt.show()
```



X.shape, y.shape

```
((48842, 14), (48842, 1))
```

Remarks of Detecting Outliers

As you can see there's 48,842 instances that have been detected.

✓ Removing Outliers using IQR

Trimming, or truncating, is the process of removing observations that show outliers in one or more variables in the dataset. There are three commonly used methods to set the boundaries beyond which a value can be considered an outlier

```

def IQR_outliers(df, column_name, thresh=1.5):
    Q1 = df[column_name].quantile(0.25)
    Q3 = df[column_name].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - (thresh * IQR)
    upper_bound = Q3 + (thresh * IQR)

    return df[(df[column_name] >= lower_bound) & (df[column_name] <= upper_bound)]

indices_to_keep = np.array([])
for col in columns:
    outliers_removed_X = IQR_outliers(X, col)
    indices_to_keep = np.intersect1d(indices_to_keep, outliers_removed_X.index)


X = X.loc[indices_to_keep]
y = y.loc[indices_to_keep]

fig, axs = plt.subplots(1, len(columns), figsize=(30, 8))

for i, col in enumerate(columns):
    value_counts = X[col].value_counts()
    axs[i].bar(value_counts.index, value_counts.values, color='skyblue')
    axs[i].set_xlabel(col)

plt.tight_layout()
plt.show()

```



`X.shape, y.shape`

```
((15221, 14), (15221, 1))
```

Verifying after Removing Outliers

Using IQR Rule the instances are now 15,221 and 27,621 instances have been removed

✓ Standardization of Data

```
from sklearn.preprocessing import StandardScaler
```

```
normalize = StandardScaler()
X_norm = normalize.fit_transform(X)
```

✓ Feature Selection

```

from sklearn.ensemble import RandomForestClassifier

clf = RandomForestClassifier(n_estimators=100, random_state=42)
clf.fit(X_norm, y)

importances = clf.feature_importances_
indices = np.argsort(importances)[::-1]


k = 4
top_k_features = indices[:k]
top_features_df = pd.DataFrame({'Feature': [f'Feature {i}' for i in top_k_features],
                                'Importance': importances[top_k_features]})

plt.figure(figsize=(10, 6))
sns.barplot(data=top_features_df, x='Feature', y='Importance', palette='virid:
plt.title(f'Top {k} Features Importance')
plt.xlabel('Feature')
plt.ylabel('Importance')
plt.show()

```

```
<ipython-input-29-5ef8c1877652>:4: DataConver
clf.fit(X_norm, y)
<ipython-input-29-5ef8c1877652>:15: FutureWar
```

Passing `palette` without assigning `hue` is



```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 15221 entries, 2 to 48839
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   15221 non-null  int64
1   workclass             15221 non-null  int64
2   fnlwgt               15221 non-null  int64
3   education             15221 non-null  int64
4   education-num        15221 non-null  int64
5   marital-status       15221 non-null  int64
6   occupation            15221 non-null  int64
7   relationship         15221 non-null  int64
8   race                 15221 non-null  int64
9   sex                  15221 non-null  int64
```

```

10 capital-gain    15221 non-null  int64
11 capital-loss    15221 non-null  int64
12 hours-per-week  15221 non-null  int64
13 native-country  15221 non-null  int64
dtypes: int64(14)
memory usage: 1.7 MB

```

✓ Remarks for Feature Selection

The result was there are 4 most important feature variable here such as `fnlwgt`, `age`, `occupation`, and `hours-per-week` which may be a factor for predicting the exact income level of an individual based on those features. This one is good for classification problems whether the person's income can actually exceeds to 50,000 or not or for regression problem to predict the exact income level of these people.

```
X = X_norm[:, [2, 0, 6, 12]]
```

```
print(X)
```

```

[[ 4.40332349e-01  8.79850060e-02 -1.54653857e-01 -4.21795495e-01]
 [ 1.24967974e+00  1.35457865e-03 -6.36041661e-01 -4.21795495e-01]
 [ 1.01945868e-01 -4.31797558e-01  8.60400447e-02 -4.21795495e-01]
 ...
 [ 7.87442068e-01 -3.45167131e-01  8.08121750e-01 -4.21795495e-01]
 [ 4.37667242e-01  1.74615433e-01  8.08121750e-01 -1.42294932e+00]
 [ 2.31103830e+00  8.79850060e-02  8.08121750e-01  2.08108908e+00]]

```

✓ Creating a Model

```

from keras.models import Sequential
from keras.layers import Input, Dense
from keras.optimizers import SGD

model = Sequential([Dense(20, activation='relu', input_shape=(X.shape[1],)),
                    Dense(20, activation='relu'),
                    Dense(1, activation='sigmoid')])

model.compile(optimizer=SGD(learning_rate=0.009),
              loss='binary_crossentropy',
              metrics=['accuracy'])

print(f"Accuracy: {model.evaluate(X, y, verbose=0)[1]*100:.2f}%")

Accuracy: 51.87%

```


✓ T1. Saving the model in HDF5 format

```
!pip install h5py
```

```
Requirement already satisfied: h5py in /usr/local/lib/python3.10/dist-packages  
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.10/dist-packages
```

```
from google.colab import drive  
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, run drive.mount('/content/drive') again
```

```
model.save("/content/drive/MyDrive/CPE018 - A8.1 / model.h5")  
print("Saved the model weights in HDF5 format to Google Drive.")
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:316: UserWarning: Saving the model weights in HDF5 format to Google Drive.  
saving_api.save_model(  
Saved the model weights in HDF5 format to Google Drive.
```

Remarks for HDF5 Format

It's efficient, versatile data storage format widely used in scientific computing and machine learning for its speed, compression, and support for hierarchical data structures.

✓ T2. Save a model and load in JSON Format

```
from tensorflow.keras.models import model_from_json
```

```
model_json = model.to_json()  
with open("model.json", "w") as json_file:  
    json_file.write(model_json)
```

```
model.save("/content/drive/MyDrive/CPE018 - A8.1 / model_json.h5")  
print("Saved the model in JSON format to Google Drive.")
```

```
Saved the model in JSON format to Google Drive.
```

```

json_file = open('model.json', 'r')
loaded_model_json = json_file.read()
json_file.close()
loaded_model = model_from_json(loaded_model_json)

loaded_model.load_weights("/content/drive/MyDrive/CPE018 - A8.1 / model_json")
print("Loaded model from disk")

Loaded model from disk

loaded_model.compile(loss='binary_crossentropy', optimizer='rmsprop', metrics=
score = loaded_model.evaluate(X, y, verbose=0)
print("%s: %.2f%%" % (loaded_model.metrics_names[1], score[1]*100))

accuracy: 51.87%

```

Remarks for JSON Format

JSON's broader compatibility and adoption make it a preferred choice for many applications.

✓ T3. Save a model and load in YAML Format


```

model_yaml = model.to_json()
with open("model.yaml", "w") as yaml_file:
    yaml_file.write(model_yaml)

model.save("/content/drive/MyDrive/CPE018 - A8.1 / model_yaml.h5")
print("Saved the model in YAML format to Google Drive.")

Saved the model in YAML format to Google Drive.
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:316
    saving_api.save_model(

```



```

yaml_file = open('model.yaml', 'r')
loaded_model_yaml = yaml_file.read()
yaml_file.close()
loaded_model = model_from_json(loaded_model_yaml)

loaded_model.load_weights("/content/drive/MyDrive/CPE018 - A8.1 / model_yaml")
print("Loaded model from disk")

Loaded model from disk

```

```
loaded_model.compile(loss='binary_crossentropy', optimizer='rmsprop', metric
score = loaded_model.evaluate(X, y, verbose=0)
print("%s: %.2f%%" % (loaded_model.metrics_names[1], score[1]*100))

accuracy: 51.87%
```

Remarks for YAML Format

YAML and JSON are both data formats used for storing and exchanging information. YAML is prized for its human-readable structure, using indentation for organization, which makes it easy to understand and write. JSON, on the other hand, is widely supported across different platforms and languages, making it popular for web development and APIs. While YAML is simpler and more intuitive for humans, JSON's broader compatibility and adoption make it a preferred choice for many applications.

✓ Verifying the loaded model from Disk

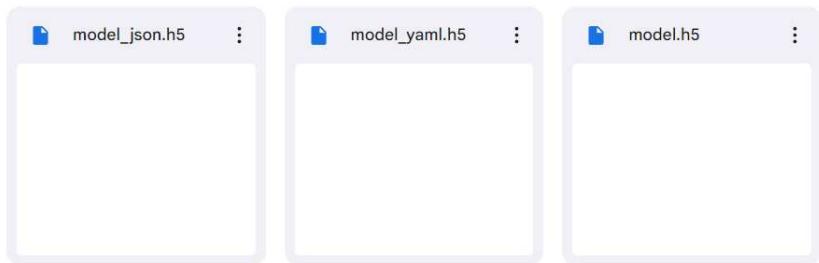
My Drive > CPE018 - A8.1 ▾

Type ▾

People ▾

Modified ▾

Files



✓ T4. Checkpoint Neural Network Model Improvements

```
from keras.callbacks import ModelCheckpoint
```

```
filepath = "weights-improvement-{epoch:02d}-{val_accuracy:.2f}.keras"
checkpoint = ModelCheckpoint(filepath, monitor='val_accuracy', verbose=1, save
```

```
callbacks_list = [checkpoint]
```

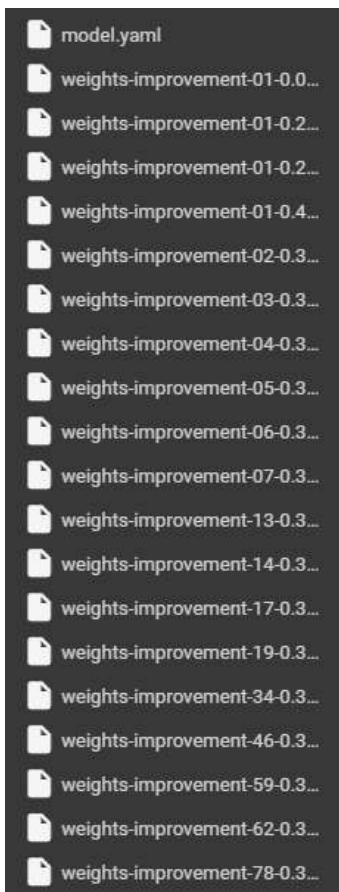
```
model.fit(X, y, validation_split=0.5, epochs=100, batch_size=20, callbacks=ca
```



Epoch 94: val_accuracy did not improve from 0.37722

Epoch 95: val_accuracy improved from 0.37722 to 0.42228, saving model

Verifying the Model Improvement



Remarks for Model Improvement

As you can see above it only save the a checkpoint if there's an improvement in model's performances.

✓ T5. Checkpoint Best Neural Network Model only

```
filepath= "weights.best.hdf5"  
checkpoint = ModelCheckpoint(filepath, monitor='val_accuracy', verbose=1, save_best_only=True, save_weights_only=True)  
callbacks_list = [checkpoint]
```

```
model.fit(X, y, validation_split=0.5, epochs=100, batch_size=20, callbacks=callbacks_list)
```



```
Epoch 93: val_accuracy did not improve from 0.44396
Epoch 94: val_accuracy did not improve from 0.44396
Epoch 95: val_accuracy did not improve from 0.44396
Epoch 96: val accuracy did not improve from 0.44396
```



Remarks

It only save the best weights that was generated in HDF5 format. This is very useful to practice experimenting the accuracy of the model.

✓ T6. Load a saved Neural Network Model

```
model = Sequential()
model.add(Dense(20, input_dim=4, kernel_initializer = 'uniform' , activation= 'relu' ))
model.add(Dense(20, kernel_initializer= 'uniform' , activation= 'relu' ))
model.add(Dense(1, kernel_initializer= 'uniform' , activation= 'sigmoid' ))

model.load_weights("weights.best.hdf5")

model.compile(loss= 'binary_crossentropy' , optimizer= 'adam' , metrics=['ac
print("Created model and loaded weights from file")
```

Created model and loaded weights from file

Remarks for Saving Neural Network Model

As you can see above, it was successfully create a model and loaded weights from the file. It demonstrate of giving an access to best weights that's being generated above.

✓ T7. Visualized Model Training History in Keras

```

from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
import matplotlib.pyplot as plt

np.random.seed(10)
X_train1, X_test1, y_train1, y_test1 = train_test_split(X, y, test_size=0.33)

model1 = Sequential([
    Dense(3, activation='relu', input_shape=(X_train1.shape[1],)),
    Dense(2, activation='relu'),
    Dense(1, activation='sigmoid')
])



optimizer = SGD(learning_rate=0.009)
model1.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['ac

history = model1.fit(X_train1, y_train1, validation_data=(X_test1, y_test1),

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```

✓ T8. Show the applicaiton of Droupout Regularization

A simple hidden layer dropout was added in the model that was created.

```
np.random.seed(10)
X_train2, X_test2, y_train2, y_test2 = train_test_split(X, y, test_size=0.5,

from tensorflow.keras.layers import Dense, Dropout

model2 = Sequential([
    Dense(50, activation='relu', input_shape=(4,)),
    Dropout(0.5),
    Dense(50, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])

optimizer = SGD(learning_rate = 0.05)
model2.compile(optimizer= optimizer,
               loss='binary_crossentropy',
               metrics=['accuracy'])

history2 = model2.fit(X_train2, y_train2, validation_data=(X_test2,y_test2),

train_loss, train_accuracy = model2.evaluate(X_train2, y_train2, verbose=0)
test_loss, test_accuracy = model2.evaluate(X_test2, y_test2, verbose=0)
print(f"Training Accuracy: {train_accuracy:.4f}, Training Loss: {train_loss:}
print(f"Test Accuracy: {test_accuracy:.4f}, Test Loss: {test_loss:.4f}")

Training Accuracy: 0.5242, Training Loss: nan
Test Accuracy: 0.5133, Test Loss: nan
```

Remarks for Dropout Regularization

Dropout in neural networks is used to prevent over-reliance on certain neurons. It randomly turns off some neurons during training, forcing the network to learn more robust features. It helps me to prevent overfitting and improves the model's ability to generalize to new, unseen data.

✓ T9. Show the application of Dropout on the visible layer

```
from keras.optimizers import SGD
from keras.layers import Dropout

np.random.seed(10)
model3 = Sequential([
    Dropout(0.8, input_shape=(4,)),
    Dense(5, activation='relu'),
    Dense(4, activation='relu'),
    Dense(1, activation='sigmoid')
])

optimizer = SGD(learning_rate = 0.05, momentum = 0.9)

model3.compile(optimizer= optimizer,
               loss='binary_crossentropy',
               metrics=['accuracy'])

history3 = model3.fit(X_train2, y_train2, validation_data=(X_test2,y_test2),

train_loss, train_accuracy = model3.evaluate(X_train2, y_train2, verbose=0)
test_loss, test_accuracy = model3.evaluate(X_test2, y_test2, verbose=0)
print(f"Training Accuracy: {train_accuracy:.4f}, Training Loss: {train_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}, Test Loss: {test_loss:.4f}")

Training Accuracy: 0.2625, Training Loss: 0.4945
Test Accuracy: 0.2562, Test Loss: 0.4289
```

Remarks for Dropout on visible layer

it encourages the model to learn more diverse patterns and reduces the risk of overfitting. This improves the model's ability to generalize to new data and enhances its overall performance.

✓ T10. Show the application of Dropout on the hidden layer

```

from keras.optimizers import SGD
from keras.layers import Dropout

np.random.seed(10)
model4 = Sequential([
    Dense(4, activation='relu', input_shape=(4,)),
    Dense(3, activation='relu'),
    Dropout(0.05),
    Dense(1, activation='sigmoid')
])

optimizer = SGD(learning_rate = 0.33, momentum = 0.9)

model4.compile(optimizer= optimizer,
               loss='binary_crossentropy',
               metrics=['accuracy'])

history4 = model4.fit(X_train2, y_train2, validation_data=(X_test2,y_test2),

train_loss, train_accuracy = model4.evaluate(X_train2, y_train2, verbose=0)
test_loss, test_accuracy = model4.evaluate(X_test2, y_test2, verbose=0)
print(f"Training Accuracy: {train_accuracy:.4f}, Training Loss: {train_loss:}
print(f"Test Accuracy: {test_accuracy:.4f}, Test Loss: {test_loss:.4f}")

Training Accuracy: 0.2625, Training Loss: 0.5976
Test Accuracy: 0.2562, Test Loss: 0.5326

```

Remarks for Dropout the hidden layer

The resulted to the comparison of the best model's performance to others wasn't good enough but it depends on the dataset you're relying on

✓ 11. Show the application of a time-based learning rate schedule

```

from tensorflow.keras.optimizers.legacy import SGD
from keras.layers import Dropout

model5 = Sequential([
    Dense(4, activation='relu', input_shape=(4,)),
    Dense(3, activation='relu'),
    Dropout(0.05),
    Dense(1, activation='sigmoid')
])

epochs = 100
learning_rate = 0.09

```

```
decay_rate = learning_rate / epochs
momentum = 0.9
optimizer = SGD(learning_rate=learning_rate, momentum=momentum, decay=decay_

model5.compile(optimizer= optimizer,
               loss='binary_crossentropy',
               metrics=['accuracy'])

history5 = model5.fit(X_train2, y_train2, validation_data=(X_test2,y_test2),

...

```

```
Epoch 91/100
238/238 - 1s - loss: 0.5522 - accuracy: 0.2625 - val_loss: 0.5131 - v
Epoch 92/100
238/238 - 1s - loss: 0.5523 - accuracy: 0.2625 - val_loss: 0.5125 - v
Epoch 93/100
238/238 - 1s - loss: 0.5525 - accuracy: 0.2625 - val_loss: 0.5128 - v
Epoch 94/100
238/238 - 1s - loss: 0.5523 - accuracy: 0.2625 - val_loss: 0.5112 - v
Epoch 95/100
238/238 - 1s - loss: 0.5525 - accuracy: 0.2625 - val_loss: 0.5121 - v
```

Remarks for time-based learning rate

The time-based learning rate schedule, also known as the decay schedule, adjusts the learning rate at each epoch automatically. This automated adjustment helps fine-tune the model efficiently without the need for manual intervention in every iteration.

✓ 12. Show the application of a drop-based learning rate schedule

```
from tensorflow.keras.optimizers.legacy import SGD
from keras.layers import Dropout
from tensorflow.keras.callbacks import LearningRateScheduler
import math
```

```
def step_decay(epoch):
    initial_lrate = 0.1
    drop = 0.5
    epochs_drop = 10.0
```